Using AWS Cloud Computing Services
Sonam Dawani

This project is to present the use of Amazon Web Services. The following Amazon Web Services are used in this project:

- S3
- AWS Glue
  - Crawler
  - ETL Job
- AWS Lambda
- AWS CloudWatch
- AWS Athena
- AWS SageMaker

The code and scripts implemented in this project are uploaded on Github.
https://github.com/SONAMDAWANI/UsingAWS_DataAnalysis

For this project simple data is used: movies.csv

movies

Film	Genre	Lead Studio	Audience score %	Profitability	Rotten Tomatoes %	Worldwide Gross	Year
Zack and Miri Make a Porno	Romance	The Weinstein Company	70	1.747541667	64	$41.94	2008
Youth in Revolt	Comedy	The Weinstein Company	52	1.09	68	$19.62	2010
You Will Meet a Tall Dark Stranger	Comedy	Independent	35	1.211818182	43	$26.66	2010
When in Rome	Comedy	Disney	44	0	15	$43.04	2010
What Happens in Vegas	Comedy	Fox	72	6.267647029	28	$219.37	2008
Water For Elephants	Drama	20th Century Fox	72	3.081421053	60	$117.09	2011
WALL-E	Animation	Disney	89	2.896019067	96	$521.28	2008
Waitress	Romance	Independent	67	11.0897415	89	$22.18	2007
Waiting For Forever	Romance	Independent	53	0.005	6	$0.03	2011
Valentine's Day	Comedy	Warner Bros.	54	4.184038462	17	$217.57	2010

Using S3 Buckets

Created bucket in S3 storage 'movies-data-bucket'

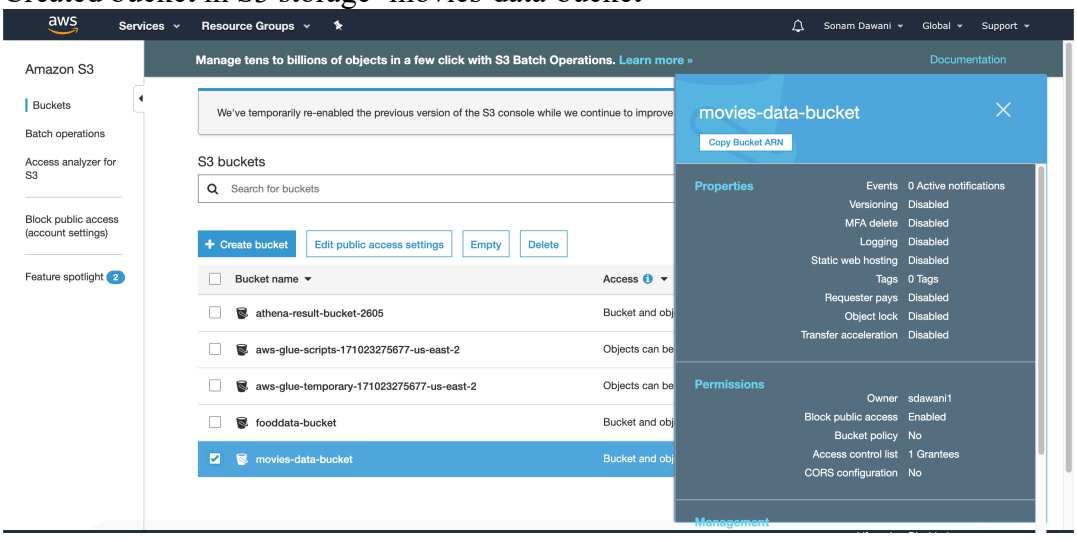

Created read and write folders in the ‘movies-data-bucket’ bucket.

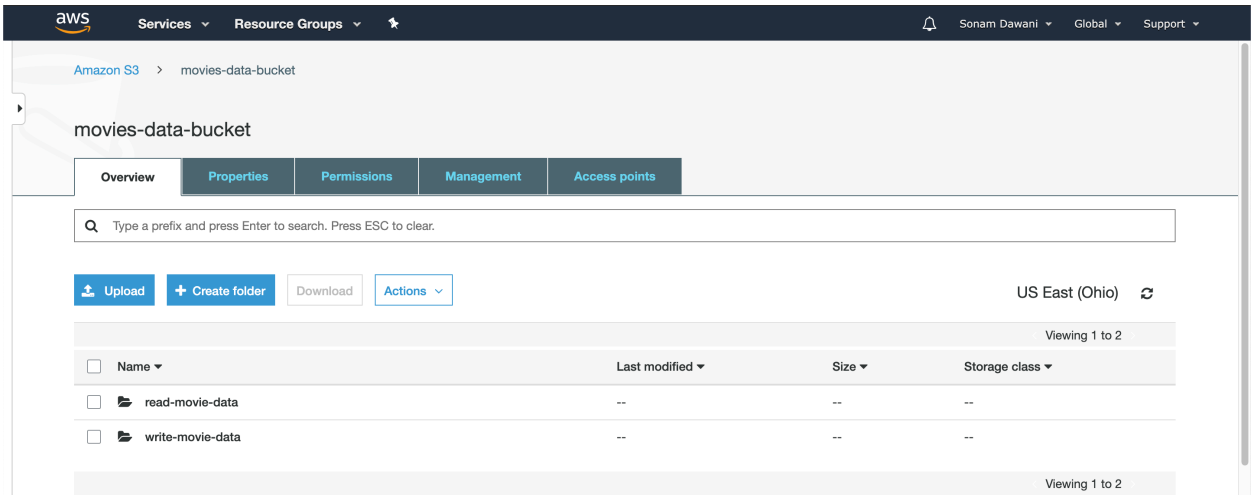

The read folder contains the original csv file, whereas write data is initially empty.
The result after ETL transformation will be stored in the write folder.

Using AWS Glue

- **Crawler**

First created a crawler which connects to the data storage (here the S3 bucket) and then creates metadata tables.

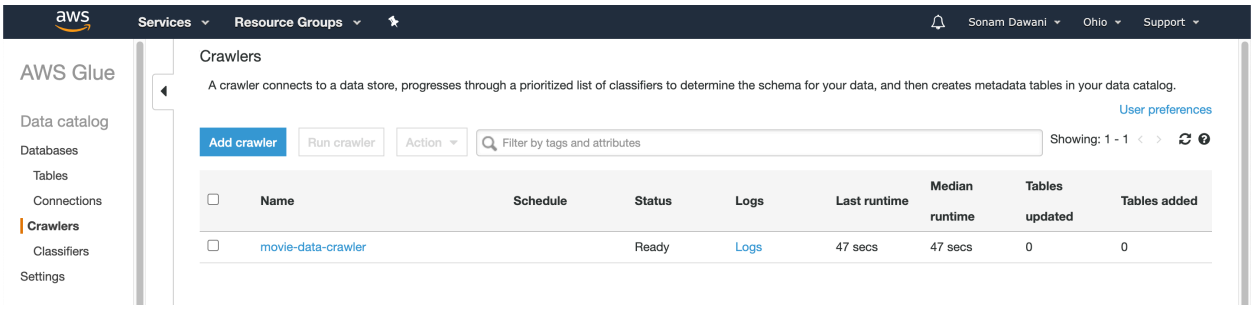

With the required IAM role and data storage information the crawler is created.

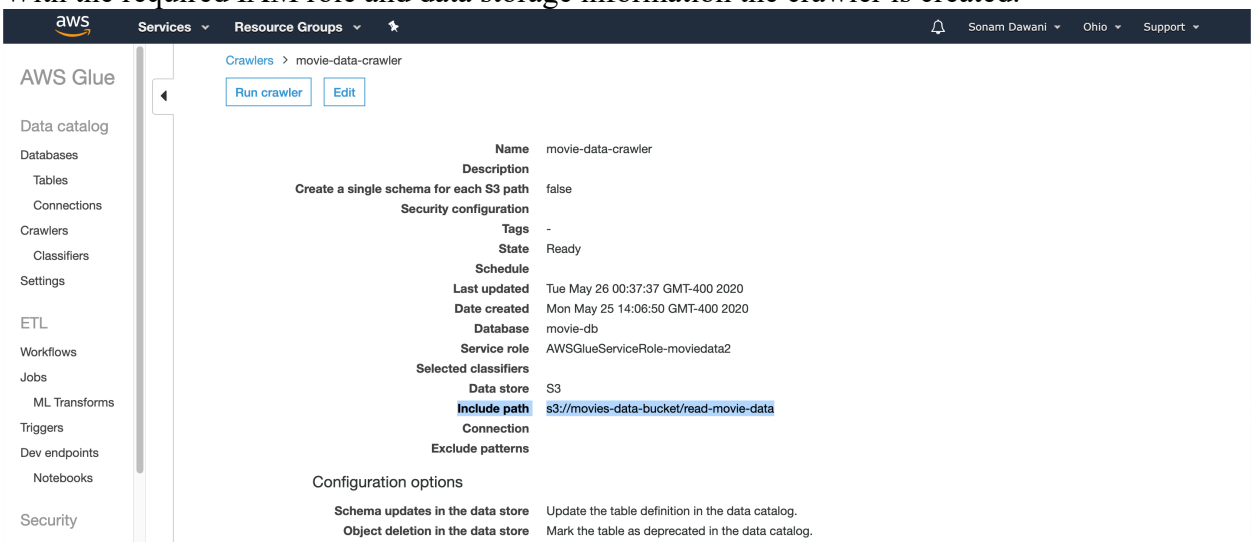

The execution of crawler above created the specified database ‘movies-db’ with ‘read_movie_data’ table in it.

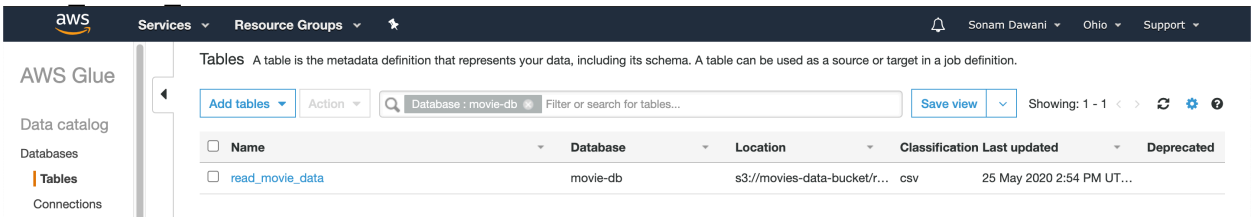

- **ETL Job**

Created run On-Demand ETL job.

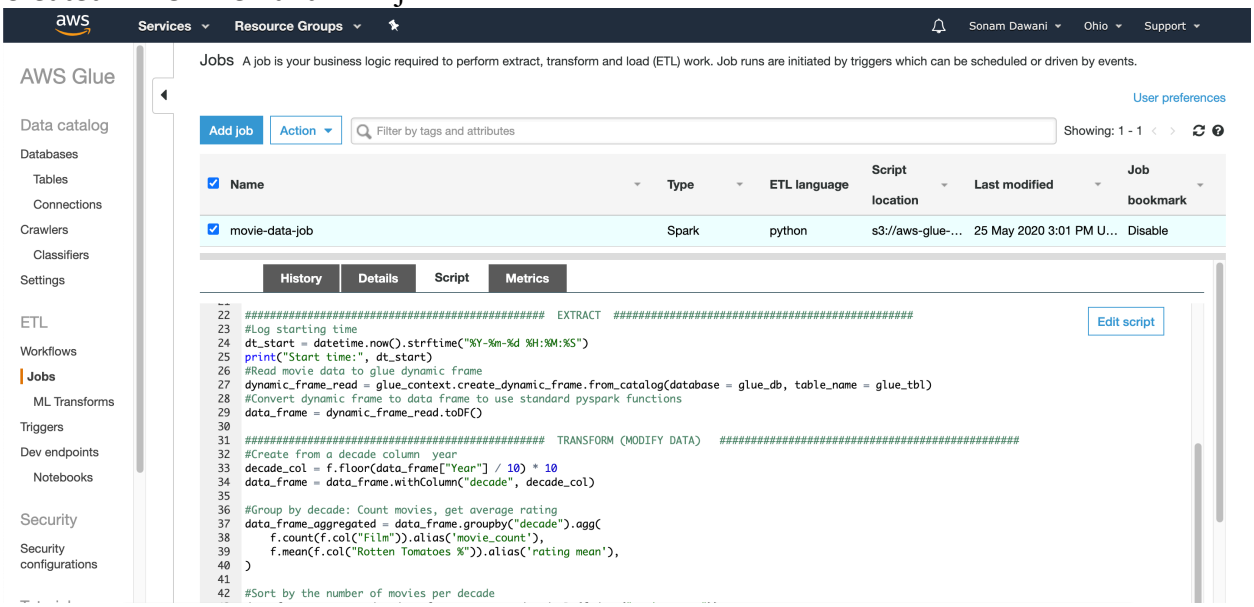

The script written for this job extract the ‘read_movie_data’ in a dataframe, performs transformation steps and load the resultant dataframe as CSV in the ‘write-movie-data’ folder of the S3 bucket.

In the transformation steps the movies are group by decades and mean rating of each decade is calculated.

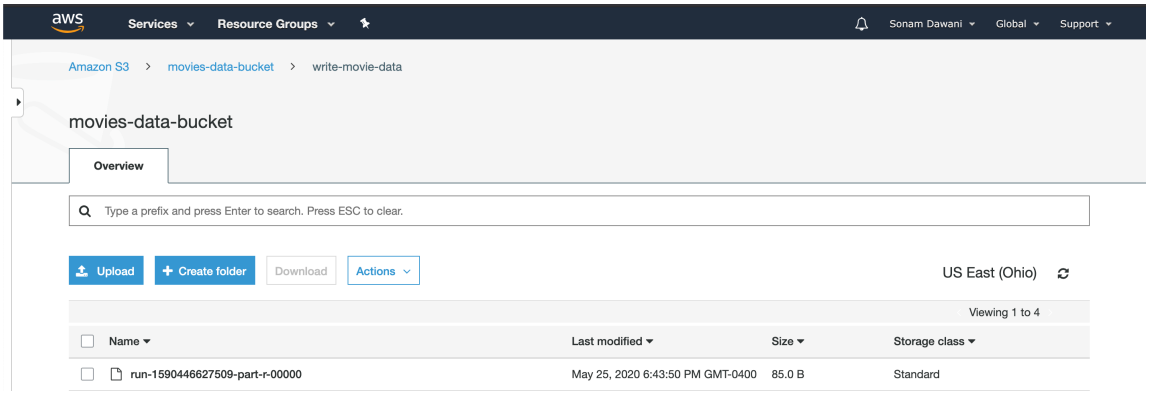

The transformed data:

run-1590446627509-part-r-00000

decade	movie_count	rating mean
2010	34	42.73529411764706
2000	43	49.83720930232558

Using AWS Lambda

AWS Lambda is used there to automate the process of running ETL job after the crawler execution is completed.

This is done in two steps:

- Creating the lambda function to start the ETL job

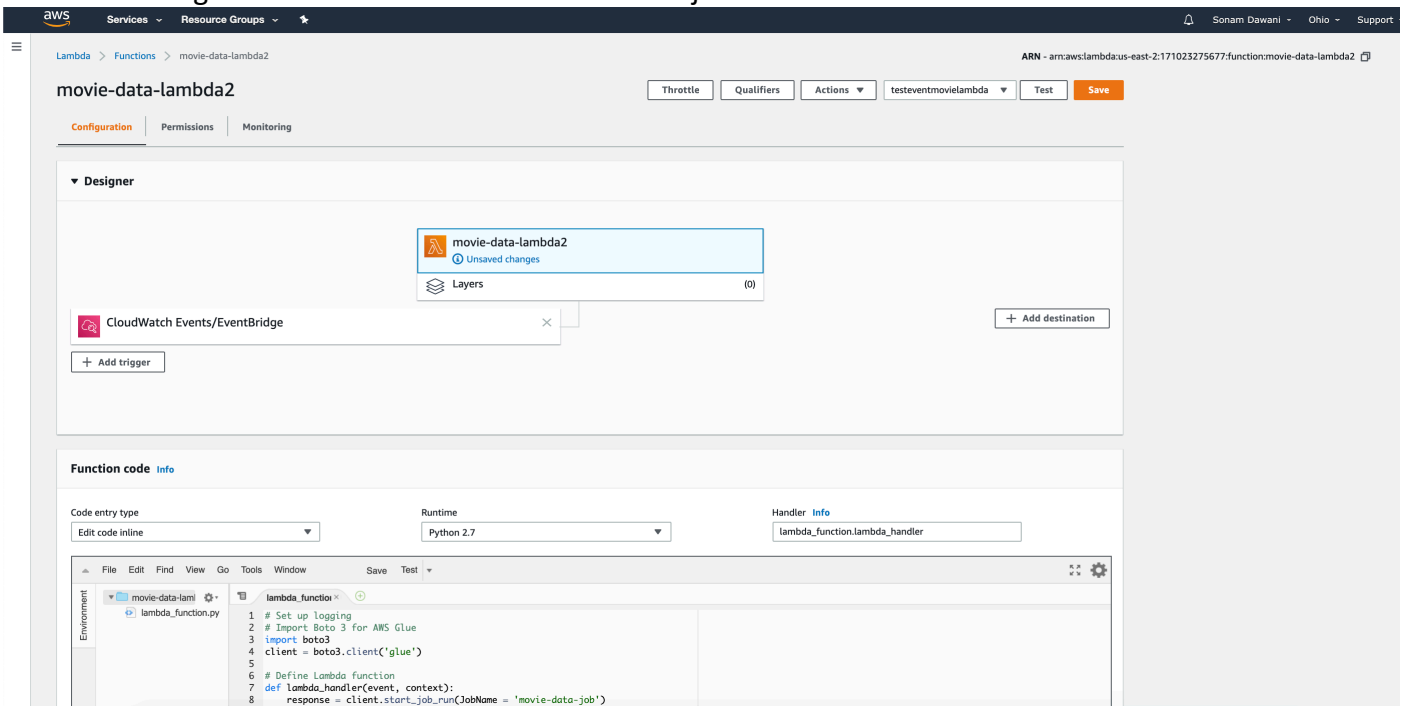

- Adding rule in CloudWatch

Created a rule in CloudWatch for the event of crawler. This rule is set to execute the lambda function once the execution of 'movie-data-crawler' is completed.

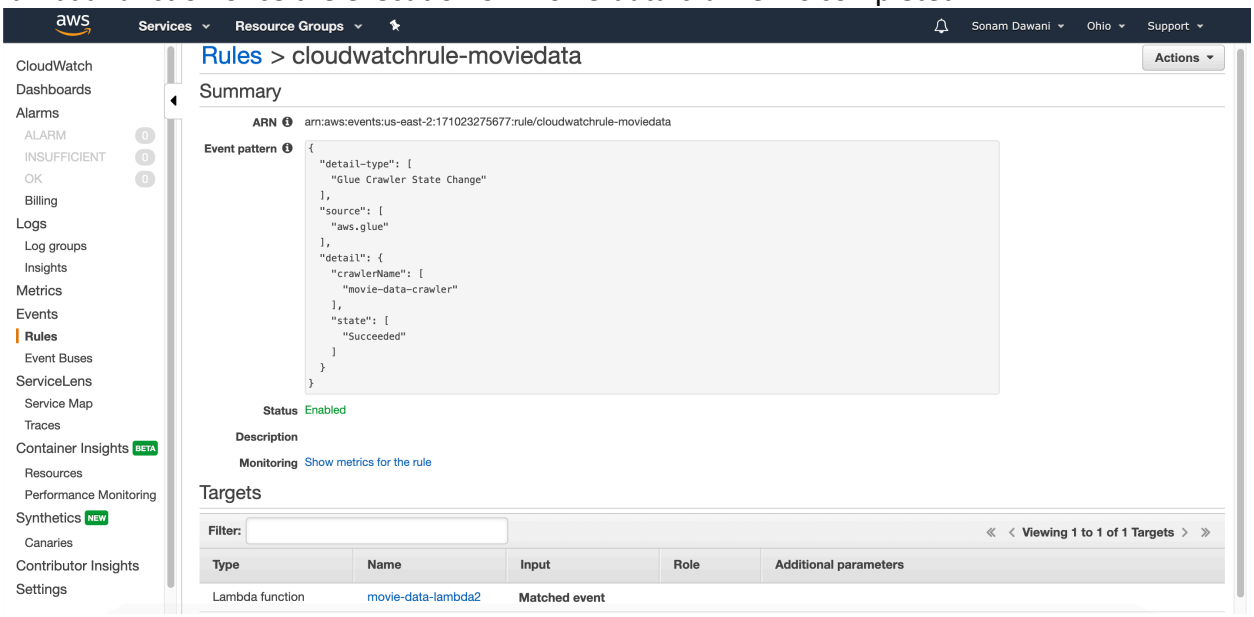

The steps above required to have **IAM roles** with some specific policies. Some of the IAM roles required to add Inline policies.

Using AWS Athena

For analysis purpose the database table was queries using AWS Athena. The results were stored in S3 bucket.

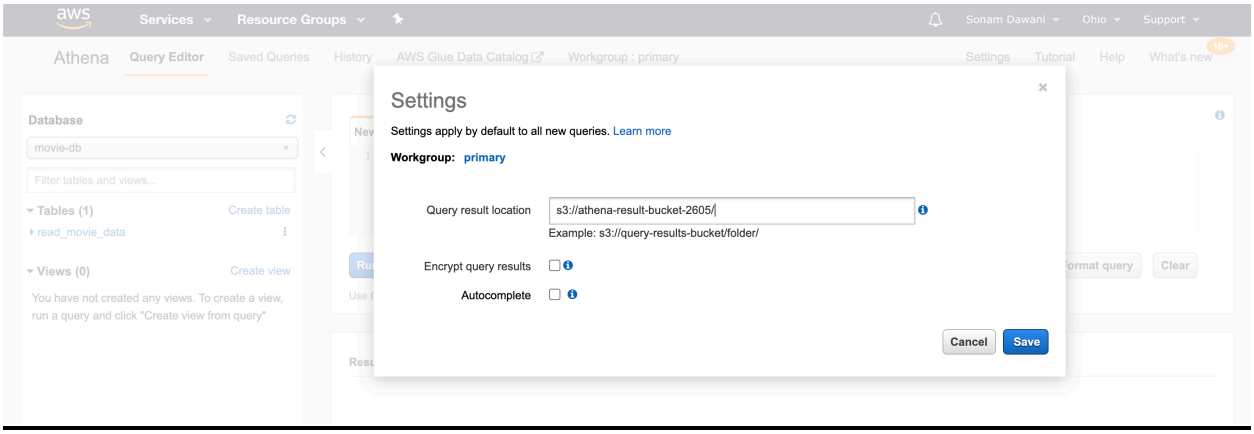

Sample of querying data:

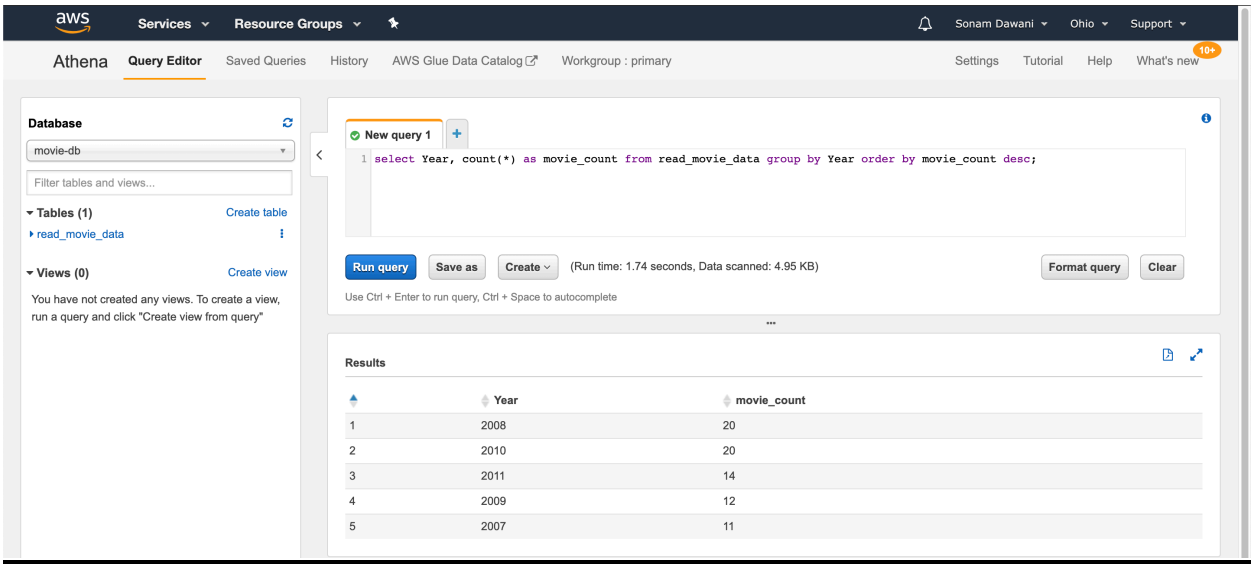

Using AWS SageMaker

For SageMaker activity I used different dataset.
The dataset used here is food data containing recipe name, time to cook, nutritional values, ingredients etc. This data is collected from Food.com

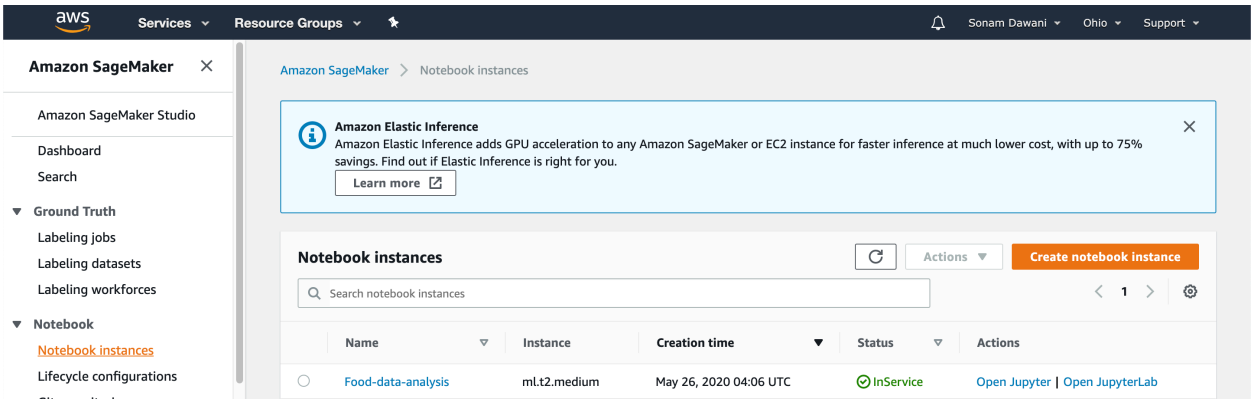

The notebook contains:

- **Data reading from S3 bucket**
- Preprocessing
- Data Exploration
- Adding cuisine data using transfer learning
- Clustering of recipes using PCA and k-means
- Regression to predict nutritional value using Gradient Boosting
- Apriori / Market Basket Analysis

The github repo contains the PDF version of the notebook.
https://github.com/SONAMDAWANI/UsingAWS_DataAnalysis